

GenAI Literacy for CS1 Students

Sondre Sæther Bolland

1 Introduction

The rise of Generative Artificial Intelligence (GenAI) has seen rapid global adoption, with the technology increasingly embedded across digital infrastructure. This is especially prevalent in computer science and programming: code editors now ship with built-in GenAI pair-programmers (GitHub Copilot, Cursor), office applications embed GenAI assistants (Microsoft Copilot, Google Gemini in Workspace), and search engines increasingly answer queries with GenAI-generated summaries rather than links to sources. For students entering computer science today, GenAI is not an optional add-on but an encompassing part of the tools they will use from their very first line of code.

While GenAI is touted by many as a step forward, a growing worry spreads through industry and academia about its effect on our own capabilities. One such worry is *deskilling*: the possibility that reliance on an AI tool causes the underlying human skill it assists with to atrophy. Studies of experts using GenAI tools in medicine [3] and software development [18] suggest that even skilled professionals can get worse at tasks their job requires as they grow more dependent on GenAI assistance. A second, related worry concerns the learning of such skills in the first place: studies of GenAI tool use in learning contexts, such as essay writing [12] and rudimentary research tasks [19], suggest these tools promote a shallower approach to learning, with learners often retaining less knowledge than peers using more traditional means, such as web search or no tool at all.

This is a particular concern for computer science education. GenAI tools have long been able to write programs expected of first-semester students in *Introductory Programming* (CS1); as early as 2022, OpenAI’s Codex model solved the large majority of typical CS1 exam and exercise questions, often outperforming the average student [10]. Four years on, these tools have improved substantially and are now freely available to anyone with an internet connection. This ability to generate all of a student’s code, paired with the risks of deskilling and shallow learning outlined above, raises substantial concerns about non-learning in CS1 for any student unable to use these tools appropriately and safely.

1.1 Scope and Aim

This paper explores this situation and provides a possible remedy. My aim lies in the question:

How can we equip CS1 students with the information needed to use GenAI without compromising their learning outcomes?

For this I turn to the concept of *AI Literacy*. Literacy is traditionally understood as the ability to identify, interpret, create, and communicate using written language [6], a concept since extended to describe comparable skill sets in other domains. *AI literacy* is one of the most recent such extensions, applying this same idea of enabling expression, communication, and access to knowledge to the domain of artificial intelligence [13].

The aim of this paper has been to develop a lecture delivered in the very first week of a CS1 course at the University of Bergen, fall 2026, helping the students use these tools safely. I start by outlining the details of AI literacy, before providing content for a lecture for AI literacy competencies. I end with a reflection on how the lecture went and how it was perceived by the students.

2 Background

In their literature review, Long and Magerko [13] synthesized a definition for AI literacy:

a set of competencies that enables individuals to critically evaluate AI technologies, communicate and collaborate effectively with AI, and use AI as a tool online, at home, and in the workplace.

This was paired with a framework of various competencies operationalizing their definition, covering five overarching themes/questions.

2.1 What is AI?

The first theme concerns the ability to recognize AI and situate it conceptually. It includes competencies such as distinguishing which technological artifacts do and do not use AI, and critically discussing what makes an entity “intelligent” at all—including how human, animal, and machine intelligence differ [13]. This theme also covers the interdisciplinary origins of AI as a field, and the distinction between *narrow* AI (systems designed to perform a single, specific task) and *general* AI (a hypothetical system capable of human-like reasoning across arbitrary domains). Much of the public discourse around “AI” conflates narrow, task-specific tools like code assistants with the far more speculative notion of general intelligence, and disentangling the two is a prerequisite for reasoning clearly about what these tools can and cannot be trusted to do [13].

2.2 What can AI do?

The second theme concerns AI’s capabilities and limitations, i.e. the ability to identify what AI is currently good and bad at. Applied to GenAI in CS1: these tools are demonstrably capable of solving the large majority of typical introductory programming exercises [10], but they are equally capable of producing code that is subtly incorrect, inefficient, insecure, or stylistically poor while still appearing fluent and confident [7]. A student who has not developed this competency has no basis for knowing when to trust a generated solution and when to interrogate it—which could lead to uncritical copy-pasting of GenAI output.

2.3 How does AI work?

The third theme is the most technically dense of the five: how AI systems represent information, make decisions, and “learn” from data, as well as the human role involved in building and training such systems [13]. For GenAI specifically, this means understanding that today’s tools are built on deep neural networks trained on vast quantities of text to predict the statistically most likely next word or code token, rather than to reason symbolically about correctness [22]. Grasping this is what allows a student to understand *why* a model can generate plausible-looking code that nonetheless fails: the system has no explicit model of program semantics,

only a learned statistical pattern over what code *typically* looks like given similar prior context.

It is worth noting, however, that more modern GenAI tools are no longer purely next-token predictors operating in isolation. Many current systems can write code, execute it in a real compiler or interpreter, observe the actual output or error message, and revise their answer accordingly, rather than relying solely on how plausible the code looks [4]. This lets a model catch at least some of the errors that pure pattern-matching would otherwise miss, though it is an empirical check bolted onto a statistical process, not a sign that the model has acquired a genuine symbolic understanding of the code it writes. Even though a script's output looks correct does not mean that the program works as intended.

2.4 How should AI be used?

The fourth theme concerns the ethical and normative dimensions of AI use, encompassing considerations such as privacy, bias, transparency, and accountability [13]. While such aspects are important, our focus lies more in how to *learn with such tools*: what constitutes appropriate use of GenAI when the explicit goal of the course is for students to develop their own programming competency, where the line falls between using a tool to learn and using it to bypass learning entirely.

2.5 How do people perceive AI?

The fifth theme concerns how humans emotionally and socially relate to AI systems, including the tendency toward anthropomorphism and an understanding that AI agents are ultimately programmable systems rather than autonomous minds [13]. As conversational GenAI tools are specifically designed to communicate in a natural, confident, human-like register, it can lead novices (such as first-semester students) to over-attribute understanding or authority to their output. A student who perceives a chatbot's fluent explanation as equivalent to an instructor's verified expertise is less likely to question it—so counteracting this perception, and re-grounding these tools as programmed systems with no ground truth of their own, is itself a literacy goal.

3 Lecture Content

I now present the content that was delivered in the first lecture of the *Introductory Programming* course, fall 2026, at the Department of Information Science and Media Studies, University of Bergen. While this material is founded on Long and Magerko's AI literacy framework [13], the lecture is limited to 60 minutes and hence the content reflects a more narrow focus, instead of all 17 competencies outlined in their paper. The lecture has been designed to address the subsections above.

Since the medium of a paper such as this differs considerably from that of a lecture, the sections that follow are not written as a speaking script, but closer to a scientific paper. I therefore omit the showmanship required to hold students' attention in the room, and instead provide additional detail elaborating on the talking points themselves. The idea is that someone could read this and develop a lecture of their own.

3.1 What is AI?

Since the inception of Artificial Intelligence as a research field, there have been many definitions of the term. Let us consider the one given by Wikipedia [1], as it is the one most readily available to a layman:

Artificial intelligence (AI) is the capability of computational systems to perform tasks typically associated with human intelligence, such as learning, reasoning, problem-solving, perception, and decision-making.

Such a definition is quite broad and puts into question to what degree such technologies embody intelligence, at least in a human sense. Let us look at some examples.

Navigation

Finding your way from one place to another was once a skill that took years to master. Ship captains relied on trained navigators who combined star charts, maps and compasses to plot a course, and even on land, route-finding across an unfamiliar city or region was a task for someone with local knowledge or a well-studied map. Today this expertise has largely been replaced by apps like Google Maps, which can compute the fastest route between two points anywhere in the world in mere seconds.

Under the hood, this feat rests on a graph problem: represent a road network as a set of points connected by roads, each labeled with a travel cost, and then find the cheapest path between two points on the graph. Dijkstra’s algorithm, published in 1959, solves exactly this problem by systematically exploring the graph outward from the starting point, always expanding the closest unvisited location next, until it has found the shortest path to every point it needs [8]. A later refinement, the A* algorithm, speeds this up by using an estimate of the remaining distance to the destination to guide the search in a more promising direction, rather than exploring equally in all directions [11].

Few people would call Google Maps intelligent, even though it performs a task—route-finding—that once required years of trained human expertise. Hence, it falls under our definition of AI. The algorithm does not understand geography, traffic, or travel in any human sense; it only follows a precise mathematical procedure over a graph that engineers built from map data.

Chess

In a 1996 six-game match, world chess champion Garry Kasparov defeated IBM’s Deep Blue—a chess bot—, winning three games to the machine’s one, with two draws. Following an upgrade, Deep Blue met Kasparov again in a 1997 rematch and reversed the result, winning two games to Kasparov’s one, with three draws—becoming the first computer to defeat a reigning world champion in a full match under tournament conditions, a result since regarded as a milestone in the history of artificial intelligence [2].

On the surface this seems to encapsulate human intelligence much more than navigation: while I myself am able to read a map and use a compass, I am far from able to beat a chess grandmaster. Yet a computer scientist would likely disagree that Deep Blue is intelligent in a human sense. The engine uses an algorithm called *minimax*: it builds out a tree of possible future moves several turns deep, scores each resulting board position with a heuristic evaluation function, and then picks the move that gives it the best guaranteed outcome, assuming its opponent always responds with their strongest possible counter-move [17]. There is no learning, no generalization, and no understanding of the game in a human sense here—only exhaustive, brute-force calculation guided by rules that human engineers wrote in advance.

GenAI

These days the term AI is used primarily to refer to GenAI tools such as Claude, ChatGPT, and Gemini—computational techniques capable of generating seemingly new, meaningful content such as text, images, or code from training data [9]. These tools have an underlying technology that can seem more complex than the previous two examples: they use *machine learning*. For us to understand how such tools are able to produce human-like text, we need to explore this underlying technique.

3.2 How does AI work?

Editorial:

Since this is the most complex part, the students were given homework before the lecture: watch 3Blue1Brown’s videos on *neural networks* and *large language models*. These give detailed, but introductory information about this technology. The lecture then proceeds to talk about these techniques more generally.

Additionally, instead of only talking about this technology while the students passively listen, they were taught the process of *induction*, i.e. the reasoning method that machine learning uses, by playing *Monkey Cards*—an unplugged computer science activity intended to teach AI to kids. We omit a description of this activity here and instead detail what was said after on *How does AI work?*

To understand how tools like Claude, ChatGPT, and Gemini can produce fluent, human-like text, we need to look at *machine learning* (ML): a branch of AI in which a system improves at a task not by following rules a human explicitly programmed—as Deep Blue’s engineers did—but by identifying statistical patterns in large amounts of data. Rather than being told *how* to solve a problem, a machine learning model is shown many examples and constructs a function that approximates the relationship between them.

Neural Networks

Most modern ML systems, including today’s GenAI tools, are built from *artificial neural networks*: layers of simple mathematical units (“neurons”) connected to one another, loosely inspired by neurons in the brain. Each connection has a numerical *weight*, and during training these weights are adjusted, incrementally and automatically, so that the network’s output gets closer to some desired result. A network with many such layers is called a *deep* neural network, which is where the term *deep learning* comes from [21]. Nobody hand-writes the rules that end up encoded in these weights; they emerge purely from repeated exposure to data.

From Neural Networks to Language

Large Language Models (LLMs)—the technology underlying GenAI chatbots—are deep neural networks built specifically to process and generate text. Since 2017, almost all state-of-the-art LLMs have used an architecture called the *transformer*, which introduced a mechanism called *self-attention*: for every word (or word-fragment, called a *token*) in a piece of text, the model computes how relevant every other token is to it, allowing it to capture relationships between words regardless of how far apart they appear in a sentence [22]. This is what allows an LLM to, for example, correctly resolve which noun a pronoun refers to several sentences earlier.

An LLM is trained by being shown enormous quantities of text scraped from the internet, books, and code repositories, with one simple task: given a sequence

of tokens, predict the next one. It is never told any grammar rules, facts, or programming syntax directly; instead, by being trained to correctly guess the next token across billions of examples, the network’s weights come to encode statistical regularities of language and code well enough to generate convincing, novel text one token at a time, each token chosen based on all tokens generated so far [22].

From text predictor to assistant

A model trained only to predict the next token is not yet a helpful assistant—left alone, it simply continues whatever text pattern it was given. To turn such a model into something like Claude or ChatGPT, developers apply an additional training stage in which human reviewers rank different model responses to the same prompt, and the model is further adjusted to produce the kinds of responses humans preferred—a technique known as *reinforcement learning from human feedback* [15]. This step is what makes the model behave like a helpful conversational partner rather than a raw autocomplete engine, though the underlying mechanism producing each word is still, fundamentally, a statistical prediction rather than an act of reasoning about truth or correctness.

Does ChatGPT Equate to Human Intelligence?

Whether a system that can fluently continue text should be considered intelligent in a human sense is a question linguist Noam Chomsky, together with Ian Roberts and Jeffrey Watumull, take up in their widely discussed 2023 essay “The False Promise of ChatGPT” [5]. Their central argument rests on a distinction between three things a system can do when confronted with a physical event: *describe* it, *predict* it, and *explain* it.

Chomsky illustrates this with the example of a dropped apple. Observing that the apple falls is a description; stating in advance that it will fall once released is a prediction. Both are useful, and both can be correct, but neither is yet an explanation. A genuine explanation, Chomsky argues, requires a further step: a counterfactual claim that *any* such object would behave the same way, tied to a causal mechanism such as gravity or the curvature of space-time. Only this last move—reasoning from a general causal principle rather than from an observed instance—constitutes what Chomsky calls thinking [5].

Applied to GenAI, the argument follows directly from how LLMs are trained, as described above: an LLM constructs a function to predict the next token by finding statistical regularities across enormous quantities of text, not by constructing an internal causal model of the world. It can therefore describe that apples fall, and predict that a given apple will fall, because both are patterns abundantly present in its training data. But Chomsky’s claim is that the model has no equivalent of the causal step: it does not *reason* that gravity is the mechanism responsible and that any comparable object would behave identically for that reason, it merely reproduces the sequence of words most statistically associated with the prompt [5]. On his account, the fluency of a model’s output can therefore mask the absence of anything resembling explanation or reasoning underneath it—a distinction directly relevant to trusting an LLM’s generated code: producing a statistically plausible solution is not the same as having reasoned about why that solution is correct.

4 Can you learn with AI tools?

Although GenAI tools do not encompass human intelligence, this does not mean they are not useful. A shortest path algorithm does not possess intelligence, yet

most people use Google Maps several times a week. Can these chatbots then, for instance, be used in a learning environment—can you learn effectively with GenAI?

4.1 AI Reliance

To explore the notion of learning with GenAI, we look to a spectrum of GenAI reliance—to what extent do you use such tools and how does that impact your learning outcomes?

Blind Usage and No Usage

The two ends of this spectrum are *blind usage*—indiscriminately copy-pasting whatever the tool generates without critically evaluating its output—and *no usage*—simply not using any such tool at all. These two extremes are fairly easy to reason about in terms of learning outcomes. Someone who does not use AI tools learns exactly as students did before such tools became commercially available in 2022. Someone who uses them blindly, on the other hand, arguably learns next to nothing.

The mere possibility of blind usage is nonetheless worth noting. That it is possible to generate an entire solution to essentially anything asked in the CS1 course [10] is important to highlight, since any student under pressure from multiple looming deadlines might be pushed toward this end of the spectrum. Being aware, in advance, that this shortcut exists—and that taking it yields no learning—is something that can help a student avoid falling into it.

4.2 AI as a Helper

The more interesting part of this AI reliance spectrum, however, is what lies between the two extremes—using AI as an aid in the learning process. This could include using such tools to research a topic, ideation, have concepts explained, get help writing text, or better understand existing code. What learning outcomes result from this kind of usage? For this, we look to studies on learning with GenAI.

Cognitive Ease at a Cost

Stadler et al. [19] investigated the cognitive load—the mental effort imposed on working memory during information processing—when students used GenAI tools compared to traditional search engines in learning tasks. Cognitive load is typically distinguished into intrinsic, extraneous, and germane load, where effective learning depends on managing these components appropriately [20]. While reducing unnecessary (extraneous) cognitive load is beneficial, sufficient engagement in germane cognitive processes—those directly related to constructing and refining knowledge—is essential for deep learning.

ChatGPT assistance was shown to reduce the effort associated with information retrieval and synthesis, thereby lowering overall cognitive load during task completion. However, this reduction appears to come at a cost. By streamlining access to ready-made explanations and solutions, GenAI tools may limit the extent to which learners actively engage in sense-making processes, such as evaluating sources, integrating information, and constructing their own understanding [20]. Hence, using such tools seems to garner a more shallow approach to the learning activity.

Your Brain on ChatGPT

Kosmyna et al.’s study *Your Brain on ChatGPT* [12] examined how access to AI tools influences writing processes. 54 adult participants were asked to write essays

within a set time frame across three sessions, each assigned throughout to one of three conditions:

- GenAI-assisted writing (using ChatGPT)
- Search engine-assisted writing (with access to Google)
- Unassisted writing (brain-only condition)

A notable finding from the study was the difficulty participants in the GenAI condition experienced in accurately recalling their own work. Immediately after finishing the first essay-writing session, participants were asked to quote a passage from what they had just written. In the first session, 83.3% of participants in the GenAI group were unable to produce an accurate quote from an essay they had completed only minutes earlier, compared to 11.1% in each of the other two groups [12]. Throughout all writing sessions, the Search Engine and Brain-only groups exhibited significantly fewer issues with quoting than those using GenAI.

Cognitive Offloading

A likely explanation for what we observe among the ChatGPT users in the studies above is *cognitive offloading*—relegation of mental effort to external tools [16]. Traditionally, such offloading has been understood as a way to redistribute effort by outsourcing narrow, well-defined sub-processes—remembering where to find information rather than the information itself, for instance, as is common with a search engine like Google [16]. GenAI tools, however, appear to introduce a more expansive form of this offloading, one that extends beyond simple information storage into the reasoning and selection processes of learning itself.

In Mittal et al.’s study [14], students pursued informal learning over the course of eight days, using either ChatGPT or Google Search to seek out information, with their information-seeking behavior tracked through a daily diary protocol. Those in the ChatGPT group were found to offload much of the work of selecting and evaluating information to the chatbot itself, resulting in a diminished sense of agency over their own learning process and, counterintuitively, a heavier meta-cognitive load than the search engine group experienced. On average, this translated into worse learning outcomes for the ChatGPT group than for the Google group [14].

5 How should the students approach CS1?

The final part of the lecture was an explanation of why I had just spent an hour talking about AI: students need to be informed of the pitfalls of these tools before starting this course, where an uncritical, over-reliant use of GenAI would quietly undermine the very skills the course exists to build.

Given these research results, it is my belief that such tools have no place in a course like *Introduction to Programming*. When you’re learning the fundamentals of a field, you have a limited ability to judge whether what is being generated is correct—and even less ability to tell whether you are actually learning material that is still largely unfamiliar to you.

I then go on to outline the official guidelines given for all assignments and the exam of the course: *no use of GenAI is allowed*. However, I state clearly that *I will conduct no plagiarism checks*. My reasoning behind this choice is to fulfill the purpose of the talk—give the students the necessary information to choose themselves. Now that they are aware of the potential issues with using such tools for learning it is up to themselves to choose whether they are going to generate their way through the course or actually do the hard thing and learn.

6 Conclusion - How did the Lecture go?

It has now been three weeks since the lecture was given. No rigorous methodology was used to collect data on students' GenAI usage or perceptions, but as course instructor, I have made informal observations throughout these weeks of lectures, seminars, and assignment hand-ins. What follows is therefore anecdotal.

During the break of the most recent lecture, three students came down to ask how to stop Visual Studio Code (the course's chosen code editor) from suggesting how to finish their lines of code—in other words, how to disable GitHub Copilot, which is built into the editor. One student stated their frustration: “I want to actually learn.” That attitude is exactly what I hoped to instill. It doesn't, however, seem to be universal.

Two hours before the first assignment deadline, I held an additional session to help any student struggling to finish the tasks or navigate our somewhat complicated hand-in system (we use *git*, a notoriously difficult system for coding novices). Among the few students who attended, I noticed several with browser tabs open to ChatGPT and Claude. One even asked which chatbot would be best for learning the course material.

A more encouraging sign came from the code submitted for the first assignment itself. While functionally correct, much of it is malformed and doesn't follow standard code quality conventions—something that would be far less likely if it had been generated by a chatbot. This wasn't true of every submission, but it was true of a great many, suggesting that a large portion of the students wrote the code themselves.

Overall, my impression is that the lecture has, to some extent, prompted students to reflect on their use of these tools—and that some of them may be learning more as a result.

References

- [1] Artificial intelligence, https://en.wikipedia.org/wiki/Artificial_intelligence, accessed: 2026-05-08
- [2] Deep blue, <https://www.ibm.com/history/deep-blue>, accessed: 2026-03-08
- [3] Budzyń, K., Romańczyk, M., Kitala, D., Kołodziej, P., Bugajski, M., Adami, H.O., Blom, J., Buszkiewicz, M., Halvorsen, N., Hassan, C., Romańczyk, T., Øyvind Holme, Jarus, K., Fielding, S., Kunar, M., Pellise, M., Pilonis, N., Kamiński, M.F., Kalager, M., Bretthauer, M., Mori, Y.: Endoscopist deskilling risk after exposure to artificial intelligence in colonoscopy: a multicentre, observational study. *The Lancet Gastroenterology Hepatology* **10**(10), 896–903 (2025). [https://doi.org/https://doi.org/10.1016/S2468-1253\(25\)00133-5](https://doi.org/https://doi.org/10.1016/S2468-1253(25)00133-5)
- [4] Chen, X., Lin, M., Schaerli, N., Zhou, D.: Teaching large language models to self-debug. In: Kim, B., Yue, Y., Chaudhuri, S., Fragkiadaki, K., Khan, M., Sun, Y. (eds.) *International Conference on Learning Representations*. vol. 2024, pp. 8746–8825 (2024), https://proceedings.iclr.cc/paper_files/paper/2024/file/2460396f2d0d421885997dd1612ac56b-Paper-Conference.pdf
- [5] Chomsky, N., Roberts, I., Watumull, J.: Noam chomsky: The false promise of chatgpt. *The New York Times* **8**(4), 177–179 (2023)
- [6] Conference, U.G.: Revised recommendation concerning the international standardization of educational statistics adopted by the General Conference at its twentieth session, Paris, 27 November 1978. Unesco (1978)

- [7] Cotroneo, D., Improta, C., Liguori, P.: Human-written vs. ai-generated code: A large-scale study of defects, vulnerabilities, and complexity. In: 2025 IEEE 36th International Symposium on Software Reliability Engineering (ISSRE). pp. 252–263 (2025). <https://doi.org/10.1109/ISSRE66568.2025.00035>
- [8] Dijkstra, E.W.: A Note on Two Problems in Connexion with Graphs, p. 287–290. Association for Computing Machinery, New York, NY, USA, 1 edn. (2022), <https://doi.org/10.1145/3544585.3544600>
- [9] Feuerriegel, S., Hartmann, J., Janiesch, C., Zschech, P.: Generative ai: S. feuerriegel et al. *Business & information systems engineering* **66**(1), 111–126 (2024)
- [10] Finnie-Ansley, J., Denny, P., Becker, B.A., Luxton-Reilly, A., Prather, J.: The robots are coming: Exploring the implications of openai codex on introductory programming. In: Proceedings of the 24th Australasian Computing Education Conference. p. 10–19. ACE '22, Association for Computing Machinery, New York, NY, USA (2022). <https://doi.org/10.1145/3511861.3511863>, <https://doi.org/10.1145/3511861.3511863>
- [11] Hart, P.E., Nilsson, N.J., Raphael, B.: A formal basis for the heuristic determination of minimum cost paths. *IEEE Transactions on Systems Science and Cybernetics* **4**(2), 100–107 (1968). <https://doi.org/10.1109/TSSC.1968.300136>
- [12] Kosmyna, N., Hauptmann, E., Yuan, Y.T., Situ, J., Liao, X.H., Beresnitzky, A.V., Braunstein, I., Maes, P.: Your brain on chatgpt: Accumulation of cognitive debt when using an ai assistant for essay writing task. *arXiv preprint arXiv:2506.08872* **4** (2025)
- [13] Long, D., Magerko, B.: What is ai literacy? competencies and design considerations. In: Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems. p. 1–16. CHI '20, Association for Computing Machinery, New York, NY, USA (2020). <https://doi.org/10.1145/3313831.3376727>
- [14] Mittal, S., Blodgett, S.L., Liao, Q.V.: Learning by chatting? investigating the impact of generative ai on information seeking and learning (2026), <https://arxiv.org/abs/2606.11669>
- [15] Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C.L., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., et al.: Training language models to follow instructions with human feedback. *arXiv preprint arXiv:2203.02155* (2022)
- [16] Risko, E.F., Gilbert, S.J.: Cognitive offloading. *Trends in cognitive sciences* **20**(9), 676–688 (2016)
- [17] Shannon, C.E.: Xxii. Programming a computer for playing chess. *The London, Edinburgh, and Dublin Philosophical Magazine and Journal of Science* **41**(314), 256–275 (1950). <https://doi.org/10.1080/14786445008521796>
- [18] Shen, J.H., Tamkin, A.: How AI impacts skill formation (2026), <https://arxiv.org/abs/2601.20245>
- [19] Stadler, M., Bannert, M., Sailer, M.: Cognitive ease at a cost: LLMs reduce mental effort but compromise depth in student scientific inquiry. *Computers in Human Behavior* **160**, 108386 (2024)

- [20] Sweller, J.: Chapter two - Cognitive Load Theory. *Psychology of Learning and Motivation*, vol. 55, pp. 37–76. Academic Press (2011). <https://doi.org/https://doi.org/10.1016/B978-0-12-387691-1.00002-8>
- [21] Sze, V., Chen, Y.H., Yang, T.J., Emer, J.S.: Efficient Processing of Deep Neural Networks: A Tutorial and Survey. *Proceedings of the IEEE* **105**(12), 2295–2329 (2017). <https://doi.org/10.1109/JPROC.2017.2761740>
- [22] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., Kaiser, L.u., Polosukhin, I.: Attention is all you need. In: Guyon, I., Luxburg, U.V., Bengio, S., Wallach, H., Fergus, R., Vishwanathan, S., Garnett, R. (eds.) *Advances in Neural Information Processing Systems*. vol. 30. Curran Associates, Inc. (2017), https://proceedings.neurips.cc/paper_files/paper/2017/file/3f5ee243547dee91fbd053c1c4a845aa-Paper.pdf